

common-task-1-electron-photon-classification

March 26, 2024

```
[ ]: pip install --upgrade --quiet tensorflow
```

Note: you may need to restart the kernel to use updated packages.

```
[ ]: import h5py
import requests
import numpy as np
import matplotlib.pyplot as plt
import os
import PIL
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.layers import Dense, Flatten, Dropout
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from sklearn.model_selection import train_test_split
from tensorflow.keras.callbacks import ModelCheckpoint
from sklearn.metrics import roc_auc_score
from tensorflow.keras.optimizers import Adam
```

```
[ ]: photon_url = "https://cernbox.cern.ch/remote.php/dav/public-files/
↳AtBT8y4MiQYFcgc/SinglePhotonPt50_IMGCRUPS_n249k_RHv1.hdf5"
photon_filename = "SinglePhotonPt50_IMGCRUPS_n249k_RHv1.hdf5"

electron_url = "https://cernbox.cern.ch/remote.php/dav/public-files/
↳FbXw3V4XNyYB3oA/SingleElectronPt50_IMGCRUPS_n249k_RHv1.hdf5"
electron_filename = "SingleElectronPt50_IMGCRUPS_n249k_RHv1.hdf5"

urls, filenames = [photon_url, electron_url], [photon_filename,
↳electron_filename]

for i in range(2):
    response = requests.get(urls[i])
    with open(filenames[i], 'wb') as f:
        f.write(response.content)
```

```
[ ]: with h5py.File(photon_filename, 'r') as f_photon:
      X_photon, y_photon = f_photon['X'][:, :], f_photon['y'][:, :]

      with h5py.File(electron_filename, 'r') as f_electron:
          X_electron, y_electron = f_electron['X'][:, :], f_electron['y'][:, :]

      X_combined = np.concatenate((X_photon, X_electron), axis=0)
      y_combined = np.concatenate((y_photon, y_electron), axis=0)

      print("Shape of combined data array:", X_combined.shape)
      print("Shape of combined target array:", y_combined.shape)
```

Shape of combined data array: (498000, 32, 32, 2)
 Shape of combined target array: (498000,)

```
[ ]: energy_channel = X_combined[:, :, :, 0]
      time_channel = X_combined[:, :, :, 1]

      # Create a new array by stacking the two channels along with a zeros array for
      # the third channel
      X_combined_3channel = np.stack([energy_channel, time_channel, np.
      zeros_like(time_channel)], axis=-1)
```

```
[ ]: x_train, x_test, y_train, y_test = train_test_split(X_combined, y_combined,
      test_size=0.2, shuffle=True, random_state=42, stratify=y_combined)
```

```
[ ]: y_train[0]
```

```
[ ]: 0.0
```

```
[ ]: x_train.shape
```

```
[ ]: (398400, 32, 32, 2)
```

```
[ ]: del X_combined_3channel
      del y_combined
      del energy_channel
      del time_channel
      del X_photon
      del y_photon
      del X_electron
      del y_electron
```

```
[ ]: model = Sequential()

      resnet50_model = keras.applications.ResNet50(
          include_top=True,
```

```

        weights=None,
        input_shape=(32, 32, 2),
        pooling='max',
        classes=1,
        classifier_activation="sigmoid",
    )
    # for layer in resnet50_model.layers:
    #     layer.trainable = False

model.add(resnet50_model)
# model.add(Flatten())
# model.add(Dense(512, activation='relu'))
# model.add(Dropout(0.5))
# model.add(Dense(1, activation='sigmoid'))

model.summary()

```

/opt/conda/lib/python3.10/site-packages/keras/src/applications/resnet.py:125:
 UserWarning: This model usually expects 1 or 3 input channels. However, it was
 passed an input_shape with 2 input channels.

```
input_shape = imagenet_utils.obtain_input_shape(
```

Model: "sequential"

Layer (type)	Output Shape	Param #
resnet50 (Functional)	?	23,586,625

Total params: 23,586,625 (89.98 MB)

Trainable params: 23,533,505 (89.77 MB)

Non-trainable params: 53,120 (207.50 KB)

```
[ ]: model.compile(loss="binary_crossentropy", optimizer=Adam(learning_rate=0.001),
    ↪metrics=["accuracy"])
```

```
[ ]: # rm -rf /kaggle/working/model_output
```

```
[ ]: output_dir = 'model_output/resnet50_maxpool_adam001'

modelcheckpoint = ModelCheckpoint(filepath=output_dir+"/weights.{epoch:02d}.
    ↪keras")
```

```

if not os.path.exists(output_dir):
    os.makedirs(output_dir)

callback = keras.callbacks.EarlyStopping(
    monitor="val_loss",
    patience=3,
    verbose=1,
    restore_best_weights=True,
)

```

```

[ ]: epochs = 50
    batch_size = 32

history = model.fit(
    x_train,
    y_train,
    batch_size=batch_size,
    epochs=epochs,
    verbose=1,
    validation_data=(x_test, y_test),
    callbacks=[callback, modelcheckpoint]
)

```

Epoch 1/50

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1711103069.426778 881 service.cc:145] XLA service 0x7f8414002570 initialized for platform CUDA (this does not guarantee that XLA will be used).
Devices:

I0000 00:00:1711103069.426822 881 service.cc:153] StreamExecutor device (0): Tesla P100-PCIE-16GB, Compute Capability 6.0

3/12450 6:48 33ms/step -
accuracy: 0.4740 - loss: 1.8538

I0000 00:00:1711103089.079931 881 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

12450/12450 416s
29ms/step - accuracy: 0.5651 - loss: 0.7258 - val_accuracy: 0.6027 - val_loss: 0.6663

Epoch 2/50

12450/12450 343s
28ms/step - accuracy: 0.6050 - loss: 0.6702 - val_accuracy: 0.6134 - val_loss: 0.6592

Epoch 3/50

12450/12450 343s

28ms/step - accuracy: 0.6331 - loss: 0.6456 - val_accuracy: 0.6649 - val_loss: 0.6316

Epoch 4/50
12450/12450 345s
28ms/step - accuracy: 0.6610 - loss: 0.6226 - val_accuracy: 0.6728 - val_loss: 0.6117

Epoch 5/50
12450/12450 341s
27ms/step - accuracy: 0.6822 - loss: 0.6028 - val_accuracy: 0.6820 - val_loss: 0.6078

Epoch 6/50
12450/12450 336s
27ms/step - accuracy: 0.6966 - loss: 0.5877 - val_accuracy: 0.6975 - val_loss: 0.5835

Epoch 7/50
12450/12450 338s
27ms/step - accuracy: 0.7056 - loss: 0.5783 - val_accuracy: 0.7092 - val_loss: 0.5720

Epoch 8/50
12450/12450 340s
27ms/step - accuracy: 0.7087 - loss: 0.5733 - val_accuracy: 0.7072 - val_loss: 0.5751

Epoch 9/50
12450/12450 340s
27ms/step - accuracy: 0.7122 - loss: 0.5686 - val_accuracy: 0.6910 - val_loss: 0.5904

Epoch 10/50
12450/12450 341s
27ms/step - accuracy: 0.7142 - loss: 0.5668 - val_accuracy: 0.7141 - val_loss: 0.5670

Epoch 11/50
12450/12450 343s
28ms/step - accuracy: 0.7171 - loss: 0.5629 - val_accuracy: 0.7129 - val_loss: 0.5670

Epoch 12/50
12450/12450 339s
27ms/step - accuracy: 0.7204 - loss: 0.5592 - val_accuracy: 0.7163 - val_loss: 0.5651

Epoch 13/50
12450/12450 359s
29ms/step - accuracy: 0.7223 - loss: 0.5560 - val_accuracy: 0.7184 - val_loss: 0.5612

Epoch 14/50
12450/12450 362s
29ms/step - accuracy: 0.7212 - loss: 0.5564 - val_accuracy: 0.7160 - val_loss: 0.5608

Epoch 15/50
12450/12450 363s

```

29ms/step - accuracy: 0.7236 - loss: 0.5534 - val_accuracy: 0.6699 - val_loss:
0.6014
Epoch 16/50
12450/12450 368s
30ms/step - accuracy: 0.7269 - loss: 0.5506 - val_accuracy: 0.7175 - val_loss:
0.5643
Epoch 17/50
12450/12450 352s
28ms/step - accuracy: 0.7268 - loss: 0.5496 - val_accuracy: 0.7212 - val_loss:
0.5569
Epoch 18/50
12450/12450 342s
27ms/step - accuracy: 0.7277 - loss: 0.5488 - val_accuracy: 0.7227 - val_loss:
0.5552
Epoch 19/50
12450/12450 342s
27ms/step - accuracy: 0.7293 - loss: 0.5466 - val_accuracy: 0.5319 - val_loss:
0.8686
Epoch 20/50
12450/12450 341s
27ms/step - accuracy: 0.7310 - loss: 0.5447 - val_accuracy: 0.7200 - val_loss:
0.5588
Epoch 21/50
12450/12450 342s
27ms/step - accuracy: 0.7324 - loss: 0.5422 - val_accuracy: 0.6662 - val_loss:
0.6207
Epoch 21: early stopping
Restoring model weights from the end of the best epoch: 18.

```

```
[ ]: model.save("trainedResNet50.keras")
```

```
[ ]: x = keras.models.load_model('trainedResNet50.keras')
```

```

/opt/conda/lib/python3.10/site-packages/keras/src/saving/saving_lib.py:396:
UserWarning: Skipping variable loading for optimizer 'rmsprop', because it has
216 variables whereas the saved optimizer has 431 variables.
  trackable.load_own_variables(weights_store.get(inner_path))

```

```
[ ]: x.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
resnet50 (Functional)	(32, 1)	23,586,625

Total params: 47,120,132 (179.75 MB)

Trainable params: 23,533,505 (89.77 MB)

Non-trainable params: 53,120 (207.50 KB)

Optimizer params: 23,533,507 (89.77 MB)

```
[ ]: import os
import subprocess
from IPython.display import FileLink, display

def download_file(path, download_file_name):
    os.chdir('/kaggle/working/')
    zip_name = f"/kaggle/working/{download_file_name}.zip"
    command = f"zip {zip_name} {path} -r"
    result = subprocess.run(command, shell=True, capture_output=True, text=True)
    if result.returncode != 0:
        print("Unable to run zip command!")
        print(result.stderr)
        return
    display(FileLink(f'{download_file_name}.zip'))
```

```
[ ]: download_file('/kaggle/working/trainedResNet50.keras', 'trainedResNet50.keras')
```

/kaggle/working/trainedResNet50.keras.zip

```
[ ]: model = Sequential()

resnet50_model = keras.applications.ResNet50(
    include_top=True,
    weights=None,
    input_shape=(32, 32, 2),
    pooling='avg',
    classes=1,
    classifier_activation="sigmoid",
)
# for layer in resnet50_model.layers:
#     layer.trainable = False

model.add(resnet50_model)
# model.add(Flatten())
# model.add(Dense(512, activation='relu'))
```

```

# model.add(Dropout(0.5))
# model.add(Dense(1, activation='sigmoid'))

model.summary()
epochs = 50
batch_size = 32

history = model.fit(
    x_train,
    y_train,
    batch_size=batch_size,
    epochs=epochs,
    verbose=1,
    validation_data=(x_test, y_test),
    callbacks=[callback, modelcheckpoint]
)

```

Epoch 1/50

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1711099400.521850 362 service.cc:145] XLA service 0x7e0c9c005fc0 initialized for platform CUDA (this does not guarantee that XLA will be used).

Devices:

I0000 00:00:1711099400.521893 362 service.cc:153] StreamExecutor device (0): Tesla P100-PCIE-16GB, Compute Capability 6.0

3/12450 6:37 32ms/step -
accuracy: 0.4549 - loss: 1.7112

I0000 00:00:1711099419.214896 362 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.

12450/12450 400s
28ms/step - accuracy: 0.5691 - loss: 0.7245 - val_accuracy: 0.5965 - val_loss: 0.6680

Epoch 2/50

12450/12450 335s
27ms/step - accuracy: 0.6134 - loss: 0.6711 - val_accuracy: 0.6056 - val_loss: 0.6691

Epoch 3/50

12450/12450 333s
27ms/step - accuracy: 0.6672 - loss: 0.6235 - val_accuracy: 0.6711 - val_loss: 0.6168

Epoch 4/50

12450/12450 333s
27ms/step - accuracy: 0.6861 - loss: 0.6033 - val_accuracy: 0.6932 - val_loss: 0.5926

Epoch 5/50


```
12450/12450          336s
27ms/step - accuracy: 0.6966 - loss: 0.5894 - val_accuracy: 0.6884 - val_loss:
0.6562
Epoch 6/50
12450/12450          337s
27ms/step - accuracy: 0.7010 - loss: 0.5837 - val_accuracy: 0.7055 - val_loss:
0.5942
Epoch 7/50
12450/12450          337s
27ms/step - accuracy: 0.7049 - loss: 0.5800 - val_accuracy: 0.6794 - val_loss:
1.0661
Epoch 7: early stopping
Restoring model weights from the end of the best epoch: 4.
```

```
[ ]: y_train.shape
```

```
[ ]: (398400,)
```

```
[ ]: x_train.shape
```

```
[ ]: (398400, 32, 32, 2)
```

```
[ ]:
```